

```
In [ ]: #1. Importe La base de datos a una base en Jupyter Notebook con pandas.
```

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [8]: base= "C:/Users/i7/Downloads/Walmart.csv"
data = pd.read_csv(base, sep=',')
```

```
In [9]: data.head()
```

```
Out[9]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106

```
In [11]: data
```

Out[11]:

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106
...
6430	45	28-09-2012	713173.95	0	64.88	3.997	192.013558	8.684
6431	45	05-10-2012	733455.07	0	64.89	3.985	192.170412	8.667
6432	45	12-10-2012	734464.36	0	54.47	4.000	192.327265	8.667
6433	45	19-10-2012	718125.53	0	56.47	3.969	192.330854	8.667
6434	45	26-10-2012	760281.43	0	58.85	3.882	192.308899	8.667

6435 rows × 8 columns

```
In [ ]: #2. Obtenga Los descriptivos resumen de La base de datos e identifique a Las variables numéricas y categóricas.
#¿Hay algo que le llame la atención?
```

```
In [10]: print(data.dtypes)
```

```
Store          int64
Date           object
Weekly_Sales    float64
Holiday_Flag    int64
Temperature     float64
Fuel_Price     float64
CPI            float64
Unemployment    float64
dtype: object
```

```
In [ ]: #La variable date se encuentra como object se puede ajustar como fecha
```

```
In [15]: data['Date'] = pd.to_datetime(data['Date'], format='%m/%d/%Y')
```

```
# Ahora la columna "Date" debería estar en formato de fecha
print(data.dtypes)
```

```
Store          int64
Date          datetime64[ns]
Weekly_Sales   float64
Holiday_Flag   int64
Temperature    float64
Fuel_Price     float64
CPI            float64
Unemployment   float64
dtype: object
```

```
In [ ]: #3. Evalúe si la base contiene datos perdidos
```

```
In [17]: data_perdidos = data.isnull().sum()
```

```
In [18]: print(data_perdidos)
```

```
Store      0
Date       0
Weekly_Sales  0
Holiday_Flag  0
Temperature  0
Fuel_Price  0
CPI         0
Unemployment  0
dtype: int64
```

```
In [19]: total_cells = data.size
total_missing = data.isnull().sum().sum()
missing_percentage = (total_missing / total_cells) * 100
print(f"Porcentaje de valores perdidos: {missing_percentage:.2f}%")
```

```
Porcentaje de valores perdidos: 0.00%
```

```
In [ ]: #No se registran datos perdidos
```

```
In [ ]: #4. Evalúe si alguna de las variables contiene datos atípicos (outliers)
```

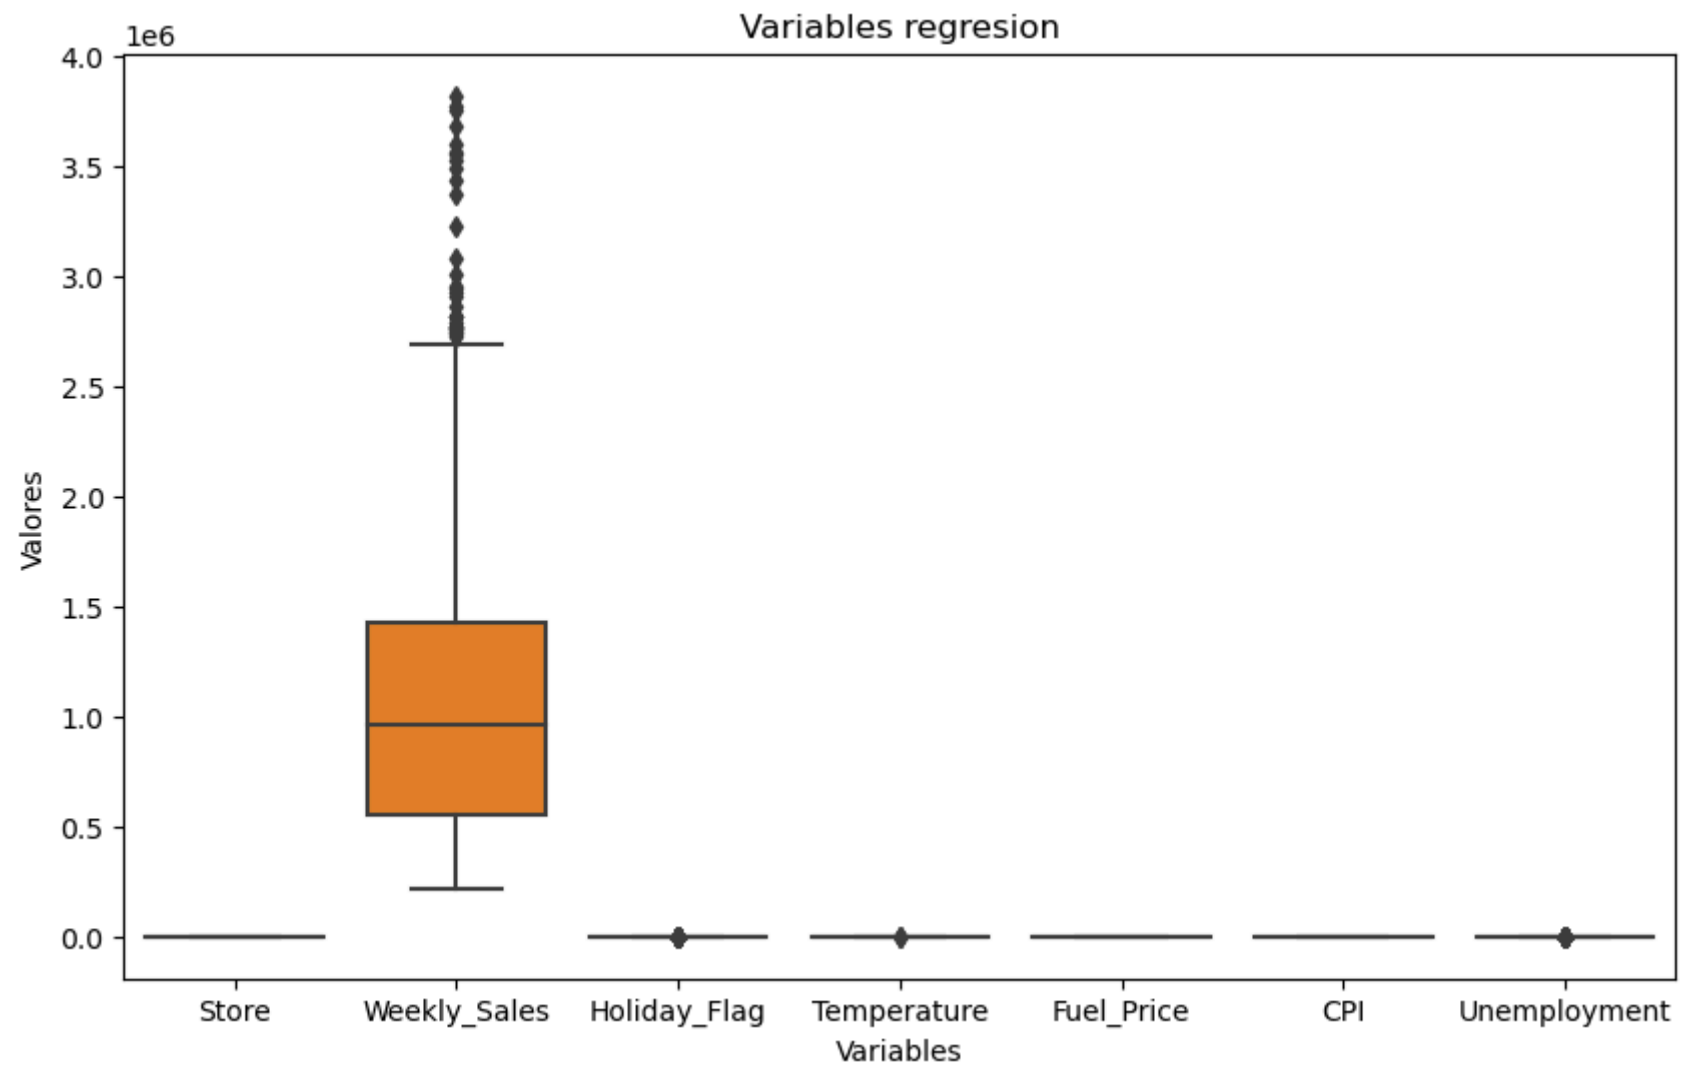
```
In [20]: data.describe()
```

Out[20]:

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
count	6435.000000	6.435000e+03	6435.000000	6435.000000	6435.000000	6435.000000	6435.000000
mean	23.000000	1.046965e+06	0.069930	60.663782	3.358607	171.578394	7.999151
std	12.988182	5.643666e+05	0.255049	18.444933	0.459020	39.356712	1.875885
min	1.000000	2.099862e+05	0.000000	-2.060000	2.472000	126.064000	3.879000
25%	12.000000	5.533501e+05	0.000000	47.460000	2.933000	131.735000	6.891000
50%	23.000000	9.607460e+05	0.000000	62.670000	3.445000	182.616521	7.874000
75%	34.000000	1.420159e+06	0.000000	74.940000	3.735000	212.743293	8.622000
max	45.000000	3.818686e+06	1.000000	100.140000	4.468000	227.232807	14.313000


In [21]:

```
plt.figure(figsize=(10, 6))
sns.boxplot(data=data)
plt.title('Variables regresion')
plt.xlabel('Variables')
plt.ylabel('Valores')
plt.show()
```



```
In [30]: data['Unemployment'].plot.box()
```

```
Out[30]: <Axes: >
```




In []: *#5.Grafique las distribuciones de las variables y a priori comente sobre ellas.*


```
In [33]: import seaborn as sns
import matplotlib.pyplot as plt

sns.displot(data['Store'], color="green", kde=True)


plt.show()
```




```
In [35]: sns.displot(data['Weekly_Sales'], color="green", kde=True)
plt.show()
```




```
In [36]: sns.displot(data['Holiday_Flag'], color="green", kde=True)
plt.show()
```

```
In [38]: sns.displot(data['Temperature'], color="green", kde=True)
plt.show()
```




```
In [39]: sns.displot(data['Fuel_Price'], color="green", kde=True)  
plt.show()
```




```
In [40]: sns.displot(data['CPI'], color="green", kde=True)

plt.show()
```




```
In [41]: sns.displot(data['Unemployment'], color="green", kde=True)
plt.show()
```




```
In [42]: sns.regplot(y='Temperature',x='Weekly_Sales',data=data,line_kws={'color':'red'})
```

```
Out[42]: <Axes: xlabel='Weekly_Sales', ylabel='Temperature'>
```



```
In [48]: sns.regplot(y='Store',x='Weekly_Sales',data=data,line_kws={'color':'red'})
```

```
Out[48]: <Axes: xlabel='Weekly_Sales', ylabel='Store'>
```



```
In [50]: data.corr(numeric_only=True).style.background_gradient(cmap='coolwarm')
```

Out[50]:

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
Store	1.000000	-0.335332	-0.000000	-0.022659	0.060023	-0.209492	0.223531
Weekly_Sales	-0.335332	1.000000	0.036891	-0.063810	0.009464	-0.072634	-0.106176
Holiday_Flag	-0.000000	0.036891	1.000000	-0.155091	-0.078347	-0.002162	0.010960
Temperature	-0.022659	-0.063810	-0.155091	1.000000	0.144982	0.176888	0.101158
Fuel_Price	0.060023	0.009464	-0.078347	0.144982	1.000000	-0.170642	-0.034684
CPI	-0.209492	-0.072634	-0.002162	0.176888	-0.170642	1.000000	-0.302020
Unemployment	0.223531	-0.106176	0.010960	0.101158	-0.034684	-0.302020	1.000000

```
In [54]: import statsmodels.api as sm
from statsmodels.formula.api import ols
regression = ols("Weekly_Sales ~ Store + Holiday_Flag + Temperature + Fuel_Price + CPI + Unemployment", data=data)
# Ajustar el modelo a los datos
results = regression.fit()
# Imprimir el resumen del modelo
print(results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Weekly_Sales    R-squared:                0.142
Model:                  OLS             Adj. R-squared:           0.141
Method:                 Least Squares   F-statistic:              176.7
Date:                   Tue, 19 Mar 2024 Prob (F-statistic):       9.33e-209
Time:                   23:33:11        Log-Likelihood:          -93861.
No. Observations:       6435           AIC:                   1.877e+05
Df Residuals:           6428           BIC:                   1.878e+05
Df Model:               6
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.996e+06	7.54e+04	26.461	0.000	1.85e+06	2.14e+06
Store	-1.539e+04	521.895	-29.486	0.000	-1.64e+04	-1.44e+04
Holiday_Flag	7.303e+04	2.59e+04	2.815	0.005	2.22e+04	1.24e+05
Temperature	-975.4019	375.974	-2.594	0.009	-1712.436	-238.367
Fuel_Price	9596.0739	1.48e+04	0.648	0.517	-1.94e+04	3.86e+04
CPI	-2319.4558	184.772	-12.553	0.000	-2681.670	-1957.241
Unemployment	-2.188e+04	3788.000	-5.776	0.000	-2.93e+04	-1.45e+04

```

=====
Omnibus:                188.961    Durbin-Watson:           0.130
Prob(Omnibus):           0.000    Jarque-Bera (JB):         205.250
Skew:                    0.435    Prob(JB):                 2.69e-45
Kurtosis:                3.100    Cond. No.                 2.19e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 2.19e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [55]: import statsmodels.api as sm
from statsmodels.formula.api import ols
```



```

regression = ols("Weekly_Sales ~ Store + Holiday_Flag + Temperature + CPI + Unemployment", data=data)
# Ajustar el modelo a los datos
results = regression.fit()
# Imprimir el resumen del modelo
print(results.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Weekly_Sales    R-squared:                0.141
Model:                  OLS             Adj. R-squared:          0.141
Method:                 Least Squares   F-statistic:             211.9
Date:                  Tue, 19 Mar 2024 Prob (F-statistic):       7.51e-210
Time:                  23:34:24         Log-Likelihood:          -93861.
No. Observations:      6435            AIC:                   1.877e+05
Df Residuals:          6429            BIC:                   1.878e+05
Df Model:               5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.032e+06	5.07e+04	40.114	0.000	1.93e+06	2.13e+06
Store	-1.537e+04	521.337	-29.488	0.000	-1.64e+04	-1.44e+04
Holiday_Flag	7.222e+04	2.59e+04	2.787	0.005	2.14e+04	1.23e+05
Temperature	-929.0252	369.081	-2.517	0.012	-1652.547	-205.503
CPI	-2345.9264	180.191	-13.019	0.000	-2699.160	-1992.693
Unemployment	-2.22e+04	3755.948	-5.910	0.000	-2.96e+04	-1.48e+04

```

=====
Omnibus:                188.685    Durbin-Watson:           0.130
Prob(Omnibus):           0.000    Jarque-Bera (JB):        204.924
Skew:                    0.434    Prob(JB):                3.17e-45
Kurtosis:                3.100    Cond. No.                1.46e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.46e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [56]: import statsmodels.api as sm
from statsmodels.formula.api import ols
regression = ols("Weekly_Sales ~ Store + Holiday_Flag + CPI + Unemployment", data=data)
# Ajustar el modelo a los datos
results = regression.fit()

```

```
# Imprimir el resumen del modelo
print(results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Weekly_Sales    R-squared:                0.141
Model:                  OLS             Adj. R-squared:          0.140
Method:                 Least Squares   F-statistic:              263.1
Date:                  Tue, 19 Mar 2024 Prob (F-statistic):       1.02e-209
Time:                  23:35:12         Log-Likelihood:           -93864.
No. Observations:      6435            AIC:                    1.877e+05
Df Residuals:          6430            BIC:                    1.878e+05
Df Model:               4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.004e+06	4.94e+04	40.528	0.000	1.91e+06	2.1e+06
Store	-1.535e+04	521.499	-29.443	0.000	-1.64e+04	-1.43e+04
Holiday_Flag	8.273e+04	2.56e+04	3.234	0.001	3.26e+04	1.33e+05
CPI	-2444.4274	175.963	-13.892	0.000	-2789.374	-2099.481
Unemployment	-2.379e+04	3703.784	-6.424	0.000	-3.11e+04	-1.65e+04

```

=====
Omnibus:                198.096    Durbin-Watson:           0.130
Prob(Omnibus):          0.000     Jarque-Bera (JB):        216.005
Skew:                   0.442     Prob(JB):                1.24e-47
Kurtosis:               3.150     Cond. No.                1.35e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 1.35e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [58]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [59]: results.predict()
```

```
Out[59]: array([1279838.7887136 , 1362216.26679386, 1279367.54050364, ...,
        636770.34491274, 636761.57235155, 636815.24048924])
```

```
In [ ]: #Se toma como referencia el enfoque econométrico, tomando los elementos de p Value < 0.05 a fin de validar la significancia
        #de las variables.
```